



Displaying Data from Multiple Tables Using Joins

Objectives

After completing this lesson, you should be able to do the following:

- Write `SELECT` statements to access data from more than one table using equijoins and nonequijoins
- Join a table to itself by using a self-join
- View data that generally does not meet a join condition by using `OUTER` joins
- Generate a Cartesian product of all rows from two or more tables

Lesson Agenda

- Types of JOINS and their syntax
 - Natural join
 - Join with the USING clause
 - Join with the ON clause
 - Self-join
 - Nonequijoins
 - OUTER join:
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
 - Cartesian product
 - Cross join


Obtaining Data from Multiple Tables

EMPLOYEES

| | EMPLOYEE_ID | FIRST_NAME | LAST_NAME | JOB_ID |
|----|-------------|------------|-----------|------------|
| 1 | 100 | Steven | King | AD_PRES |
| 2 | 101 | Neena | Kochhar | AD_VP |
| 3 | 102 | Lex | De Haan | AD_VP |
| 4 | 103 | Alexander | Hunold | IT_PROG |
| 5 | 104 | Bruce | Ernst | IT_PROG |
| 6 | 105 | David | Austin | IT_PROG |
| 7 | 106 | Valli | Pataballa | IT_PROG |
| 8 | 107 | Diana | Lorentz | IT_PROG |
| 9 | 108 | Nancy | Greenberg | FI_MGR |
| 10 | 109 | Daniel | Faviet | FI_ACCOUNT |

JOBS

| | JOB_ID | JOB_TITLE |
|---|------------|-------------------------------|
| 1 | AD_PRES | President |
| 2 | AD_VP | Administration Vice President |
| 3 | AD_ASST | Administration Assistant |
| 4 | FI_MGR | Finance Manager |
| 5 | FI_ACCOUNT | Accountant |
| 6 | AC_MGR | Accounting Manager |
| 7 | AC_ACCOUNT | Public Accountant |
| 8 | SA_MAN | Sales Manager |
| 9 | SA_REP | Sales Representative |



| | EMPLOYEE_ID | JOB_ID | JOB_TITLE |
|---|-------------|------------|-------------------------------|
| 1 | 206 | AC_ACCOUNT | Public Accountant |
| 2 | 205 | AC_MGR | Accounting Manager |
| 3 | 200 | AD_ASST | Administration Assistant |
| 4 | 100 | AD_PRES | President |
| 5 | 101 | AD_VP | Administration Vice President |
| 6 | 102 | AD_VP | Administration Vice President |
| 7 | 109 | FI_ACCOUNT | Accountant |

...

Types of Joins

Joins that are compliant with the SQL:1999 standard include the following:

- Natural join with the `NATURAL JOIN` clause
- Join with the `USING` clause
- Join with the `ON` clause
- OUTER joins:
 - `LEFT OUTER JOIN`
 - `RIGHT OUTER JOIN`
 - `FULL OUTER JOIN`
- Cross joins

Joining Tables Using SQL:1999 Syntax

Use a join to query data from more than one table:

```
SELECT    table1.column, table2.column
FROM      table1
[NATURAL JOIN table2] |
[JOIN table2 USING (column_name)] |
[JOIN table2 ON (table1.column_name = table2.column_name)] |
[LEFT|RIGHT|FULL OUTER JOIN table2
  ON (table1.column_name = table2.column_name)] |
[CROSS JOIN table2];
```

Lesson Agenda

- Types of JOINS and their syntax
- Natural join
- Join with the USING clause
- Join with the ON clause
- Self-join
- Nonequijoins
- OUTER join:
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
- Cartesian product
 - Cross join

Creating Natural Joins

- The `NATURAL JOIN` clause is based on all the columns that have the same name in two tables.
- It selects rows from the two tables that have equal values in all matched columns.
- If the columns having the same names have different data types, an error is returned.

```
SELECT * FROM table1 NATURAL JOIN table2;
```


Retrieving Records with Natural Joins

```
SELECT employee_id, first_name, job_id, job_title
from employees NATURAL JOIN jobs;
```

| | EMPLOYEE_ID | FIRST_NAME | JOB_ID | JOB_TITLE |
|----|-------------|------------|------------|-------------------------------|
| 1 | 100 | Steven | AD_PRES | President |
| 2 | 101 | Neena | AD_VP | Administration Vice President |
| 3 | 102 | Lex | AD_VP | Administration Vice President |
| 4 | 103 | Alexander | IT_PROG | Programmer |
| 5 | 104 | Bruce | IT_PROG | Programmer |
| 6 | 105 | David | IT_PROG | Programmer |
| 7 | 106 | Valli | IT_PROG | Programmer |
| 8 | 107 | Diana | IT_PROG | Programmer |
| 9 | 108 | Nancy | FI_MGR | Finance Manager |
| 10 | 109 | Daniel | FI_ACCOUNT | Accountant |
| 11 | 110 | John | FI_ACCOUNT | Accountant |



...

Creating Joins with the USING Clause

- If several columns have the same names but the data types do not match, use the `USING` clause to specify the columns for the equijoin.
- Use the `USING` clause to match only one column when more than one column matches.

Joining Column Names



EMPLOYEES

| |  EMPLOYEE_ID |  DEPARTMENT_ID |
|----|---|---|
| 1 | 200 | 10 |
| 2 | 201 | 20 |
| 3 | 202 | 20 |
| 4 | 205 | 110 |
| 5 | 206 | 110 |
| 6 | 100 | 90 |
| 7 | 101 | 90 |
| 8 | 102 | 90 |
| 9 | 103 | 60 |
| 10 | 104 | 60 |

...

Foreign key

DEPARTMENTS

| |  DEPARTMENT_ID |  DEPARTMENT_NAME |
|---|---|---|
| 1 | 10 | Administration |
| 2 | 20 | Marketing |
| 3 | 50 | Shipping |
| 4 | 60 | IT |
| 5 | 80 | Sales |
| 6 | 90 | Executive |
| 7 | 110 | Accounting |
| 8 | 190 | Contracting |

Primary key

Retrieving Records with the USING Clause

```
SELECT employee_id, last_name,  
       location_id, department_id  
FROM   employees JOIN departments  
       USING (department_id) ;
```

| | <small>A Z</small> EMPLOYEE_ID | <small>A Z</small> LAST_NAME | <small>A Z</small> LOCATION_ID | <small>A Z</small> DEPARTMENT_ID |
|---|--------------------------------|------------------------------|--------------------------------|----------------------------------|
| 1 | 200 | Whalen | 1700 | 10 |
| 2 | 201 | Hartstein | 1800 | 20 |
| 3 | 202 | Fay | 1800 | 20 |
| 4 | 144 | Vargas | 1500 | 50 |
| 5 | 143 | Matos | 1500 | 50 |
| 6 | 142 | Davies | 1500 | 50 |
| 7 | 141 | Rajs | 1500 | 50 |
| 8 | 124 | Mourgos | 1500 | 50 |

...

| | | | | |
|----|-----|---------|------|-----|
| 18 | 206 | Gietz | 1700 | 110 |
| 19 | 205 | Higgins | 1700 | 110 |

Qualifying Ambiguous Column Names

- Use table prefixes to qualify column names that are in multiple tables.
- Use table prefixes to increase the speed of parsing of the statement.
- Instead of full table name prefixes, use table aliases.
- Table alias gives a table a shorter name:
 - Keeps SQL code smaller, uses less memory
- Use column aliases to distinguish columns that have identical names, but reside in different tables.

Using Table Aliases with the USING Clause

- Do not qualify a column that is used in the NATURAL join or a join with a USING clause.
- If the same column is used elsewhere in the SQL statement, do not alias it.

```
SELECT l.city, d.department_name  
FROM   locations l JOIN departments d  
USING (location_id)  
WHERE d.location_id = 1400;
```

```
ORA-25154: column part of USING clause cannot have qualifier  
25154. 00000 - "column part of USING clause cannot have qualifier"  
*Cause:   Columns that are used for a named-join (either a NATURAL join  
          or a join with a USING clause) cannot have an explicit qualifier.  
*Action:  Remove the qualifier.  
Error at Line: 4 Column: 6
```

Creating Joins with the ON Clause

- The join condition for the natural join is basically an equijoin of all columns with the same name.
- Use the ON clause to specify arbitrary conditions or specify columns to join.
- The join condition is separated from other search conditions.
- The ON clause makes code easy to understand.

Retrieving Records with the ON Clause

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
FROM   employees e JOIN departments d  
ON     (e.department_id = d.department_id);
```

| | EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID_1 | LOCATION_ID |
|----|-------------|-----------|---------------|-----------------|-------------|
| 1 | 200 | Whalen | 10 | 10 | 1700 |
| 2 | 201 | Hartstein | 20 | 20 | 1800 |
| 3 | 202 | Fay | 20 | 20 | 1800 |
| 4 | 124 | Mourgos | 50 | 50 | 1500 |
| 5 | 144 | Vargas | 50 | 50 | 1500 |
| 6 | 143 | Matos | 50 | 50 | 1500 |
| 7 | 142 | Davies | 50 | 50 | 1500 |
| 8 | 141 | Rajs | 50 | 50 | 1500 |
| 9 | 107 | Lorentz | 60 | 60 | 1400 |
| 10 | 104 | Ernst | 60 | 60 | 1400 |
| 11 | 103 | Hunold | 60 | 60 | 1400 |

...

Creating Three-Way Joins

```
SELECT employee_id, city, department_name
FROM   employees e
JOIN   departments d
ON     d.department_id = e.department_id
JOIN   locations l
ON     d.location_id = l.location_id;
```

| |  EMPLOYEE_ID |  CITY |  DEPARTMENT_NAME |
|---|---|--|---|
| 1 | 100 | Seattle | Executive |
| 2 | 101 | Seattle | Executive |
| 3 | 102 | Seattle | Executive |
| 4 | 103 | Southlake | IT |
| 5 | 104 | Southlake | IT |
| 6 | 107 | Southlake | IT |
| 7 | 124 | South San Francisco | Shipping |
| 8 | 141 | South San Francisco | Shipping |
| 9 | 142 | South San Francisco | Shipping |

...

Applying Additional Conditions to a Join

Use the **AND** clause or the **WHERE** clause to apply additional conditions:

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
FROM   employees e JOIN departments d  
ON     (e.department_id = d.department_id)  
AND    e.manager_id = 149 ;
```

Or

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
FROM   employees e JOIN departments d  
ON     (e.department_id = d.department_id)  
WHERE  e.manager_id = 149 ;
```

Lesson Agenda

- Types of JOINS and their syntax
- Natural join
- Join with the USING clause
- Join with the ON clause
- **Self-join**
- Nonequijoins
- OUTER join:
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
- Cartesian product
 - Cross join

Joining a Table to Itself

EMPLOYEES (WORKER)

| EMPLOYEE_ID | LAST_NAME | MANAGER_ID |
|-------------|-----------|------------|
| 200 | Whalen | 101 |
| 201 | Hartstein | 100 |
| 202 | Fay | 201 |
| 205 | Higgins | 101 |
| 206 | Gietz | 205 |
| 100 | King | (null) |
| 101 | Kochhar | 100 |
| 102 | De Haan | 100 |
| 103 | Hunold | 102 |
| 104 | Ernst | 103 |

...

EMPLOYEES (MANAGER)

| EMPLOYEE_ID | LAST_NAME |
|-------------|-----------|
| 200 | Whalen |
| 201 | Hartstein |
| 202 | Fay |
| 205 | Higgins |
| 206 | Gietz |
| 100 | King |
| 101 | Kochhar |
| 102 | De Haan |
| 103 | Hunold |
| 104 | Ernst |

...



MANAGER_ID in the WORKER table is equal to
EMPLOYEE_ID in the MANAGER table.

Self-Joins Using the ON Clause

```
SELECT worker.last_name emp, manager.last_name mgr
FROM   employees worker JOIN employees manager
ON     (worker.manager_id = manager.employee_id);
```

| |  EMP |  MGR |
|---|---|---|
| 1 | Hunold | De Haan |
| 2 | Fay | Hartstein |
| 3 | Gietz | Higgins |
| 4 | Lorentz | Hunold |
| 5 | Ernst | Hunold |
| 6 | Zlotkey | King |
| 7 | Mourgos | King |
| 8 | Kochhar | King |

...

Lesson Agenda

- Types of JOINS and their syntax
- Natural join
- Join with the USING clause
- Join with the ON clause
- Self-join
- **Nonequijoins**
- OUTER join:
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
- Cartesian product
 - Cross join

Nonequijoins

EMPLOYEES

| | LAST_NAME | SALARY |
|-----|-----------|--------|
| 1 | Whalen | 4400 |
| 2 | Hartstein | 13000 |
| 3 | Fay | 6000 |
| 4 | Higgins | 12000 |
| 5 | Gietz | 8300 |
| 6 | King | 24000 |
| 7 | Kochhar | 17000 |
| 8 | De Haan | 17000 |
| 9 | Hunold | 9000 |
| 10 | Ernst | 6000 |
| ... | | |
| 19 | Taylor | 8600 |
| 20 | Grant | 7000 |

JOB_GRADES

| | GRADE_LEVEL | LOWEST_SAL | HIGHEST_SAL |
|---|-------------|------------|-------------|
| 1 | A | 1000 | 2999 |
| 2 | B | 3000 | 5999 |
| 3 | C | 6000 | 9999 |
| 4 | D | 10000 | 14999 |
| 5 | E | 15000 | 24999 |
| 6 | F | 25000 | 40000 |

The JOB_GRADES table defines the LOWEST_SAL and HIGHEST_SAL range of values for each GRADE_LEVEL. Therefore, the GRADE_LEVEL column can be used to assign grades to each employee.

Retrieving Records with Nonequijoins

```
SELECT e.last_name, e.salary, j.grade_level
FROM   employees e JOIN job_grades j
ON     e.salary
      BETWEEN j.lowest_sal AND j.highest_sal;
```

| | LAST_NAME | SALARY | GRADE_LEVEL |
|----|-----------|--------|-------------|
| 1 | Vargas | 2500 | A |
| 2 | Matos | 2600 | A |
| 3 | Davies | 3100 | B |
| 4 | Rajs | 3500 | B |
| 5 | Lorentz | 4200 | B |
| 6 | Whalen | 4400 | B |
| 7 | Mourgos | 5800 | B |
| 8 | Ernst | 6000 | C |
| 9 | Fay | 6000 | C |
| 10 | Grant | 7000 | C |

...

Lesson Agenda

- Types of JOINS and their syntax
- Natural join
- Join with the USING clause
- Join with the ON clause
- Self-join
- Nonequijoin
- **OUTER join:**
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
- Cartesian product
 - Cross join

Returning Records with No Direct Match Using OUTER Joins

DEPARTMENTS

| | DEPARTMENT_NAME | DEPARTMENT_ID |
|---|-----------------|---------------|
| 1 | Administration | 10 |
| 2 | Marketing | 20 |
| 3 | Shipping | 50 |
| 4 | IT | 60 |
| 5 | Sales | 80 |
| 6 | Executive | 90 |
| 7 | Accounting | 110 |
| 8 | Contracting | 190 |

There are no employees in department 190.

Employee "Grant" has not been assigned a department ID.

Equijoin with EMPLOYEES

| | DEPARTMENT_ID | LAST_NAME |
|----|---------------|-----------|
| 1 | 10 | Whalen |
| 2 | 20 | Hartstein |
| 3 | 20 | Fay |
| 4 | 110 | Higgins |
| 5 | 110 | Gietz |
| 6 | 90 | King |
| 7 | 90 | Kochhar |
| 8 | 90 | De Haan |
| 9 | 60 | Hunold |
| 10 | 60 | Ernst |

...




| | | |
|----|----|--------|
| 18 | 80 | Abel |
| 19 | 80 | Taylor |

INNER Versus OUTER Joins

- In SQL:1999, the join of two tables returning only matched rows is called an `INNER` join.
- A join between two tables that returns the results of the `INNER` join as well as the unmatched rows from the left (or right) table is called a left (or right) `OUTER` join.
- A join between two tables that returns the results of an `INNER` join as well as the results of a left and right join is a full `OUTER` join.

LEFT OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name
FROM   employees e LEFT OUTER JOIN departments d
ON     (e.department_id = d.department_id) ;
```

| |  LAST_NAME |  DEPARTMENT_ID |  DEPARTMENT_NAME |
|---|---|---|---|
| 1 | Whalen | 10 | Administration |
| 2 | Fay | 20 | Marketing |
| 3 | Hartstein | 20 | Marketing |
| 4 | Vargas | 50 | Shipping |
| 5 | Matos | 50 | Shipping |

...

| | | | |
|----|---------|--------|------------|
| 16 | Kochhar | 90 | Executive |
| 17 | King | 90 | Executive |
| 18 | Gietz | 110 | Accounting |
| 19 | Higgins | 110 | Accounting |
| 20 | Grant | (null) | (null) |

RIGHT OUTER JOIN

```
SELECT e.last_name, d.department_id, d.department_name
FROM   employees e RIGHT OUTER JOIN departments d
ON     (e.department_id = d.department_id) ;
```

| | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|---|-----------|---------------|-----------------|
| 1 | Whalen | 10 | Administration |
| 2 | Hartstein | 20 | Marketing |
| 3 | Fay | 20 | Marketing |
| 4 | Davies | 50 | Shipping |
| 5 | Vargas | 50 | Shipping |
| 6 | Rajs | 50 | Shipping |
| 7 | Mourgos | 50 | Shipping |
| 8 | Matos | 50 | Shipping |

...

| | | | |
|----|---------|-----|-------------|
| 18 | Higgins | 110 | Accounting |
| 19 | Gietz | 110 | Accounting |
| 20 | (null) | 190 | Contracting |

FULL OUTER JOIN

```
SELECT e.last_name, d.department_id, d.department_name
FROM   employees e FULL OUTER JOIN departments d
ON     (e.department_id = d.department_id) ;
```

| | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|---|-----------|---------------|-----------------|
| 1 | King | 90 | Executive |
| 2 | Kochhar | 90 | Executive |
| 3 | De Haan | 90 | Executive |
| 4 | Hunold | 60 | IT |

...

| | | | |
|----|-----------|--------|----------------|
| 15 | Grant | (null) | (null) |
| 16 | Whalen | 10 | Administration |
| 17 | Hartstein | 20 | Marketing |
| 18 | Fay | 20 | Marketing |
| 19 | Higgins | 110 | Accounting |
| 20 | Gietz | 110 | Accounting |
| 21 | (null) | 190 | Contracting |

Lesson Agenda

- Types of JOINS and their syntax
- Natural join
- Join with the USING clause
- Join with the ON clause
- Self-join
- Nonequijoins
- OUTER join:
 - LEFT OUTER join
 - RIGHT OUTER join
 - FULL OUTER join
- Cartesian product
 - Cross join

Cartesian Products

- Cartesian product is a join of every row of one table to every row of another table.
- A Cartesian product generates a large number of rows and the result is rarely useful.

Generating a Cartesian Product

EMPLOYEES (20 rows)

| | EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID |
|-----|-------------|-----------|---------------|
| 1 | 200 | Whalen | 10 |
| 2 | 201 | Hartstein | 20 |
| 3 | 202 | Fay | 20 |
| 4 | 205 | Higgins | 110 |
| ... | | | |
| 19 | 176 | Taylor | 80 |
| 20 | 178 | Grant | (null) |

DEPARTMENTS (8 rows)

| | DEPARTMENT_ID | DEPARTMENT_NAME | LOCATION_ID |
|---|---------------|-----------------|-------------|
| 1 | 10 | Administration | 1700 |
| 2 | 20 | Marketing | 1800 |
| 3 | 50 | Shipping | 1500 |
| 4 | 60 | IT | 1400 |
| 5 | 80 | Sales | 2500 |
| 6 | 90 | Executive | 1700 |
| 7 | 110 | Accounting | 1700 |
| 8 | 190 | Contracting | 1700 |




Cartesian product:
20 x 8 = 160 rows

| | EMPLOYEE_ID | DEPARTMENT_ID | LOCATION_ID |
|-----|-------------|---------------|-------------|
| 1 | 200 | 10 | 1700 |
| 2 | 201 | 20 | 1700 |
| ... | | | |
| 21 | 200 | 10 | 1800 |
| 22 | 201 | 20 | 1800 |
| ... | | | |
| 159 | 176 | 80 | 1700 |
| 160 | 178 | (null) | 1700 |

Creating Cross Joins

- A CROSS JOIN is a JOIN operation that produces the Cartesian product of two tables.
- To create a Cartesian product, specify the CROSS JOIN in your SELECT statement.

```
SELECT last_name, department_name  
FROM employees  
CROSS JOIN departments ;
```

| |  LAST_NAME |  DEPARTMENT_NAME |
|---|---|---|
| 1 | Abel | Administration |
| 2 | Davies | Administration |
| 3 | De Haan | Administration |
| 4 | Ernst | Administration |
| 5 | Fay | Administration |

...

| | | |
|-----|---------|-------------|
| 158 | Vargas | Contracting |
| 159 | Whalen | Contracting |
| 160 | Zlotkey | Contracting |

Quiz

If you join a table to itself, what kind of join are you using?

- a. Nonequijoin
- b. Left OUTER join
- c. Right OUTER join
- d. Full OUTER join
- e. Self-join
- f. Natural join
- g. Cartesian products

Summary

In this lesson, you should have learned how to :

- Write `SELECT` statements to access data from more than one table using equijoins and nonequijoins
- Join a table to itself by using a self-join
- View data that generally does not meet a join condition by using `OUTER` joins
- Generate a Cartesian product of all rows from two or more tables

Practice 7: Overview

This practice covers the following topics:

- Joining tables using an equijoin
- Performing outer and self-joins
- Adding conditions